


Selected memory improves feasibility, but not the maximum reliably

Within-run comparison: 264 parents, with a no-memory control generated from the same parent.



B. Equal-budget scientific tail
Mean best-of-13 reward difference versus no memory



Interpretation
Memory clearly raises feasibility; both tail intervals cross zero.
So this run does not show a reliable average gain in the maximum.

Bootstrap intervals resample matched parent groups within this single run; they are not across-run confidence intervals.